

CMSC 471
Intro to AI
Search (Cont.)

General search algorithm

;; problem describes the start state, operators, goal test, and operator costs
;; queueing-function is a comparator function that ranks two states
;; general-search returns either a goal node or failure

```
function general-search (problem, QUEUEING-FUNCTION)


  nodes = MAKE-QUEUE (MAKE-NODE (problem.INITIAL-STATE))
  loop
    if EMPTY(nodes) then return "failure"
    node = REMOVE-FRONT(nodes)
    if problem.GOAL-TEST (node.STATE) succeeds
      then return node
    nodes = QUEUEING-FUNCTION (nodes, EXPAND (node,
      problem.OPERATORS))
  end
```

;; Note: The goal test is NOT done when nodes are generated
;; Note: This algorithm does not detect loops

Properties of Searching Strategies


- **Completeness**
 - Guarantees finding a solution whenever one exists
- **Time complexity** (worst or average case)
 - Usually measured by *number of nodes expanded*
- **Space complexity**
 - Usually measured by maximum size of graph/tree during the search
- **Optimality/Admissibility**
 - If a solution is found, is it **guaranteed** to be an optimal one, i.e., one with minimum cost

Uninformed Search (Recap)



Search Strategy	Expanded nodes
BFS	
DFS	
Iterative-Deepening Search	
Uniform-Cost Search	

Uninformed Search (Recap)



Search Strategy	Expanded nodes
BFS	SABCDEG
DFS	SADEG
Iterative-Deepening Search	SSABCSEADEG
Uniform-Cost Search	SADBCEG

Comparing Search Strategies

Search Strategy	Time	Space	Complete?	Optimal?
BFS				
DFS				
Iterative-Deepening Search				
Uniform-Cost Search				

b = Branching Factor ; d = depth

Comparing Search Strategies

Search Strategy	Time	Space	Complete?	Optimal?
BFS	b^d			
DFS	b^d			
Iterative-Deepening Search	b^d			
Uniform-Cost Search	b^d			

b = Branching Factor ; d = depth

Comparing Search Strategies

Search Strategy	Time	Space	Complete?	Optimal?
BFS	b^d	b^d		
DFS	b^d	bd		
Iterative-Deepening Search	b^d	bd		
Uniform-Cost Search	b^d	b^d		

b = Branching Factor ; d = depth

Comparing Search Strategies

Search Strategy	Time	Space	Complete?	Optimal?
BFS	b^d	b^d	Yes	
DFS	b^d	bd	No	
Iterative-Deepening Search	b^d	bd	Yes	
Uniform-Cost Search	b^d	b^d	Yes	

b = Branching Factor ; d = depth

Comparing Search Strategies

Search Strategy	Time	Space	Complete?	Optimal?
BFS	b^d	b^d	Yes	Yes (Only for unitary cost)
DFS	b^d	bd	No	No
Iterative-Deepening Search	b^d	bd	Yes	Yes (Only for unitary cost)
Uniform-Cost Search	b^d	b^d	Yes	Yes

b = Branching Factor ; d = depth

Notes on Uniform-Cost Search

- Greedy Algorithm
 - Algorithms that make locally optimal choice in the hopes of reaching global optima
- Uniform because:
 - Tries to reach uniform cost on the priority queue or node list

UCS vs BFS vs DFS

- DFS:
 - Goes too far down the depth before backtracking
 - Can find solution in less time
- BFS:
 - Goes too far down the width before backtracking
 - Optimal
- UCS
 - Tries to find a balance between going too far deep or wide

Informed (Heuristic) Search

- Heuristic search
- Best-first search
 - Greedy search
 - Beam search
 - A* Search
- Heuristic functions

Big idea: [heuristic](#)

Merriam-Webster's Online Dictionary:

Heuristic (pron. \hyu-'ris-tik\): adj. [from Greek *heuriskein* to discover] involving or serving as an aid to learning, discovery, or problem-solving by experimental and especially trial-and-error methods

Heuristics, More Formally

$h(n)$ is a **heuristic function**, that maps a state n
to an estimated cost from n -to-goal

Heuristics, More Formally

$h(n)$ is a **heuristic function**, that maps a state n to an estimated cost from n -to-goal

$h(n)$ is **admissible** iff $h(n) \leq$ the lowest actual cost from n -to-goal

Heuristics, More Formally

$h(n)$ is a **heuristic function**, that maps a state n to an estimated cost from n -to-goal

$h(n)$ is **admissible** iff $h(n) \leq$ the lowest actual cost from n -to-goal

$h(n)$ is **consistent** iff
 $h(n) \leq \text{lowestcost}(n, n') + h(n')$

Informed methods add domain-specific information

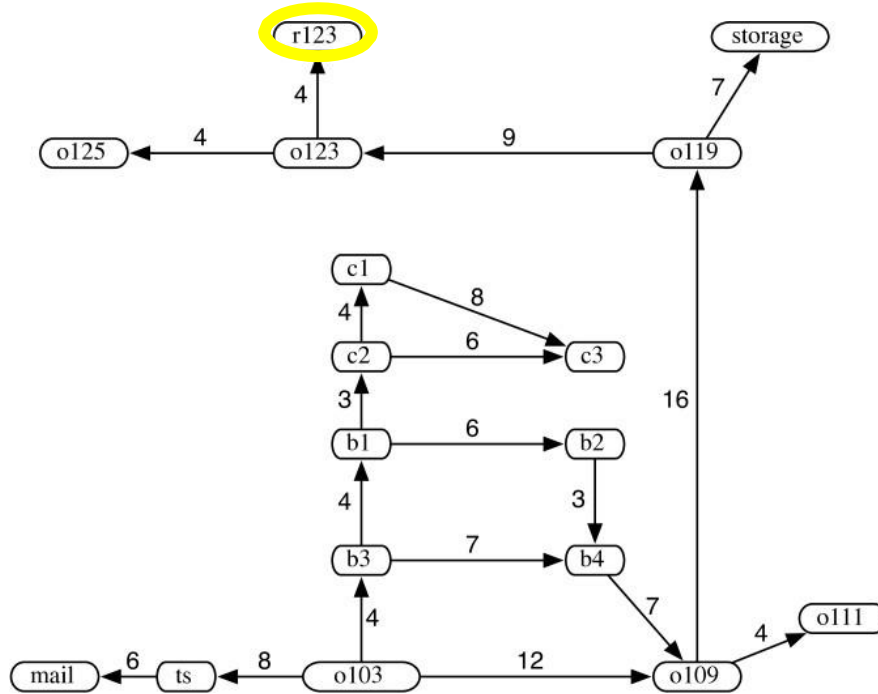
- Select best path along which to continue searching
- $h(n)$: estimates *goodness* of node n
- $h(n)$ = **estimated cost** (or distance) of minimal cost path from n **to a goal state**.
- Based on domain-specific information and computable from current state description that estimates how close we are to a goal

Heuristics

- **All domain knowledge** used in search is encoded in the **heuristic function, $h(<node>)$**
- Examples:
 - 8-puzzle: number of tiles out of place
 - 8-puzzle: sum of distances each tile is from its goal

In general


- $h(n) \geq 0$ for all nodes n
- $h(n) = 0$ implies that n is a goal node
- $h(n) = \infty$ implies n is a dead-end that can't lead to goal



Example 3.5

$$h(o123) = 4 \quad h(o125) = 6 \quad h(r123) = 0$$


(Partial) Heuristic $h(n)$
for goal **r123**



(Partial) Heuristic $h(n)$
for goal **r123**

Example 3.5

$h(mail) = 26$	$h(ts) = 23$	$h(o103) = 21$
$h(o109) = 24$	$h(o111) = 27$	$h(o119) = 11$
$h(o123) = 4$	$h(o125) = 6$	$h(r123) = 0$
$h(b1) = 13$	$h(b2) = 15$	$h(b3) = 17$
$h(b4) = 18$	$h(c1) = 6$	$h(c2) = 10$
$h(c3) = 12$	$h(storage) = 12$	




(Partial) Heuristic $h(n)$
for goal **r123**

Example 3.5

Q: Is this an **admissible** heuristic?

$h(mail) = 26$	$h(ts) = 23$	$h(o103) = 21$
$h(o109) = 24$	$h(o111) = 27$	$h(o119) = 11$
$h(o123) = 4$	$h(o125) = 6$	$h(r123) = 0$
$h(b1) = 13$	$h(b2) = 15$	$h(b3) = 17$
$h(b4) = 18$	$h(c1) = 6$	$h(c2) = 10$
$h(c3) = 12$	$h(storage) = 12$	



(Partial) Heuristic $h(n)$
for goal **r123**

Example 3.5

Q: Is this an **admissible** heuristic?

Q: Is it an accurate heuristic?

$h(mail) = 26$	$h(ts) = 23$	$h(o103) = 21$
$h(o109) = 24$	$h(o111) = 27$	$h(o119) = 11$
$h(o123) = 4$	$h(o125) = 6$	$h(r123) = 0$
$h(b1) = 13$	$h(b2) = 15$	$h(b3) = 17$
$h(b4) = 18$	$h(c1) = 6$	$h(c2) = 10$
$h(c3) = 12$	$h(storage) = 12$	

Heuristics for 8-puzzle

The number of misplaced tiles (not including the blank)

Current State

1	2	3
4	5	6
7		8

Goal State

1	2	3
4	5	6
7	8	

1	2	3
4	5	6
7		8

N	N	N
N	N	N
N	Y	

In this case, only “8” is misplaced, so heuristic function evaluates to 1

In other words, the heuristic *says* that it *thinks* a solution may be available in just 1 more move

Heuristics for 8-puzzle

Manhattan Distance (not including the blank)

Current State

3	2	8
4	5	6
7	1	

Goal State

1	2	3
4	5	6
7	8	

- The **3**, **8** and **1** tiles are misplaced (by 2, 3, and 3 steps) so the heuristic function evaluates to 8
- Heuristic says that it *thinks* a solution may be available in just 8 more moves.
- The misplaced heuristic's value is 3

3	→	<u>3</u>

2 spaces

	←	8
	↓	
	<u>8</u>	

3 spaces


<u>1</u>	←	
	↑	
	1	

3 spaces

Total 8

We can use heuristics
to guide search

Manhattan Distance
heuristic helps us
quickly find a
solution to the 8-
puzzle



Best-first search

- Search algorithm that improves **depth- first search** by expanding most promising node chosen according to heuristic rule
- Order nodes on nodes list by increasing value of an evaluation function, **$f(n)$** , incorporating domain-specific information

Best-first search


- Search algorithm that improves **depth- first search** by expanding most promising node chosen according to heuristic rule
- Order nodes on nodes list by increasing value of an evaluation function, $f(n)$, incorporating domain-specific information
- This is a generic way of referring to the class of informed methods

Greedy best first search

- A greedy algorithm makes locally optimal choices in hope of finding a global optimum
- Uses evaluation function $f(n) = h(n)$, sorting nodes by increasing values of f
- Selects node to expand appearing **closest** to goal (i.e., node with smallest f value)
- Not complete
- Not Admissible

Greedy best first search example

- Proof of non-admissibility
 - Assume arc costs = 1, greedy search finds goal g , with solution cost of 5
 - Optimal solution is path to goal with cost 3



Beam search

- Use evaluation function $f(n)$, but maximum size of the nodes list is k , a fixed constant
- Only keep k best nodes as candidates for expansion, discard rest
- k is the *beam width*
- More space efficient than greedy search, but may discard nodes on a solution path
- As k increases, approaches best first search
- Complete?
- Admissible?

Beam search

- Use evaluation function $f(n)$, but maximum size of the nodes list is k , a fixed constant
- Only keep k best nodes as candidates for expansion, discard rest
- k is the *beam width*
- More space efficient than greedy search, but may discard nodes on a solution path
- As k increases, approaches best first search
- **Not Complete**
- **Not Admissible**

We've *got* to be able to do
better, right?

A* Search

Use an evaluation function

$$f(n) = g(n) + h(n)$$

estimated **total cost** from
start to goal via state n



minimal-cost path from
the start state to state n



cost estimate from state n
to the goal

A* Search

- Use an evaluation function

$$f(n) = g(n) + h(n)$$

estimated total cost from start to goal via state n  minimal-cost path from the start state to state n  cost estimate from state n to the goal

- $g(n)$ term adds “breadth-first” component to evaluation function
- Ranks nodes on search frontier by estimated cost of solution from start node *via given node* to goal

A*

- Pronounced “*a star*”
- h is **admissible** when $h(n) \leq h^*(n)$ holds
 - $h^*(n)$ = *true cost of minimal cost path* from n to a goal
- Using an admissible heuristic guarantees that 1st solution found will be an **optimal** one
- A* is **complete** whenever branching factor is finite and every action has fixed, positive cost
- A* is **admissible**

Implementing A*

Q: Can this be an instance of our general search algorithm?




Figure 3.3

Implementing A*

Q: Can this be an instance of our general search algorithm?

A: Yup! Just make the fringe a priority queue ordered by $f(n)$




Figure 3.3

Alternative A* Pseudo-code

- 1 Put the start node S on the nodes list, called OPEN
- 2 If OPEN is empty, exit with failure
- 3 Select node in OPEN with minimal $f(n)$ and place on CLOSED
- 4 If n is a goal node, collect path back to start and stop
- 5 Expand n , generating all its successors and attach to them pointers back to n . For each successor n' of n
 - 1 If n' not already on OPEN or CLOSED
 - put n' on OPEN
 - compute $h(n')$, $g(n')=g(n)+c(n,n')$, $f(n')=g(n')+h(n')$
 - 2 If n' already on OPEN or CLOSED and if $g(n')$ is lower for new version of n' , then:
 - Redirect pointers backward from n' on path with lower $g(n')$
 - Put n' on OPEN

Observations on A^*

- **Perfect heuristic:** If $h(n) = h^*(n)$ for all n , only nodes on an optimal solution path expanded; no extra work is done

Observations on A^*

- **Perfect heuristic:** If $h(n) = h^*(n)$ for all n , only nodes on an optimal solution path expanded; no extra work is done
- **Null heuristic:** If $h(n) = 0$ for all n , then it is an admissible heuristic and A^* acts like uniform-cost search

Observations on A^*

- **Perfect heuristic:** If $h(n) = h^*(n)$ for all n , only nodes on an optimal solution path expanded; no extra work is done
- **Null heuristic:** If $h(n) = 0$ for all n , then it is an admissible heuristic and A^* acts like uniform-cost search
- **Better heuristic:** If $h_1(n) < h_2(n) \leq h^*(n)$ for all non-goal nodes, then h_2 is a *better* heuristic than h_1

Observations on A^*

- **Perfect heuristic:** If $h(n) = h^*(n)$ for all n , only nodes on an optimal solution path expanded; no extra work is done
- **Null heuristic:** If $h(n) = 0$ for all n , then it is an admissible heuristic and A^* acts like uniform-cost search
- **Better heuristic:** If $h_1(n) < h_2(n) \leq h^*(n)$ for all non-goal nodes, then h_2 is a *better* heuristic than h_1
 - If A_1^* uses h_1 , and A_2^* uses h_2 , then every node expanded by A_2^* is also expanded by A_1^*
 - i.e., A_1 expands at least as many nodes as A_2^*
 - We say that A_2^* is *better informed* than A_1^*

Observations on A^*

- **Perfect heuristic:** If $h(n) = h^*(n)$ for all n , only nodes on an optimal solution path expanded; no extra work is done
- **Null heuristic:** If $h(n) = 0$ for all n , then it is an admissible heuristic and A^* acts like uniform-cost search
- **Better heuristic:** If $h_1(n) < h_2(n) \leq h^*(n)$ for all non-goal nodes, then h_2 is a *better* heuristic than h_1
 - If A_1^* uses h_1 , and A_2^* uses h_2 , then every node expanded by A_2^* is also expanded by A_1^*
 - i.e., A_1 expands at least as many nodes as A_2^*
 - We say that A_2^* is *better informed* than A_1^*
- The closer h to h^* , the fewer extra nodes expanded

Proof of the optimality of A^*

- Assume that A^* has selected $G2$, a goal state with a suboptimal solution, i.e., $g(G2) > f^*$
- Proof by contradiction shows it's impossible

Proof of the optimality of A^*

- Assume that A^* has selected $G2$, a goal state with a suboptimal solution, i.e., $g(G2) > f^*$
- Proof by contradiction shows it's impossible
 - Choose a node n on an optimal path to G
 - Because $h(n)$ is admissible, $f^* \geq f(n)$
 - If we choose $G2$ instead of n for expansion, then $f(n) \geq f(G2)$
 - This implies $f^* \geq f(G2)$
 - $G2$ is a goal state: $h(G2) = 0$, $f(G2) = g(G2)$.
 - Therefore $f^* \geq g(G2)$
 - Contradiction

How to find good heuristics

Some options (mix-and-match):

- If $h_1(n) < h_2(n) \leq h^*(n)$ for all n , h_2 is better than (**dominates**) h_1
- **Relaxing problem:** remove constraints for easier problem; use its solution cost as heuristic function
- Max of two admissible heuristics is a **Combining heuristics**: admissible heuristic, and it's better!
- Use statistical estimates to compute h ; may lose admissibility
- Identify good features, then use **machine learning** to find heuristic function; also may lose admissibility