

CMSC 471: Intro to AI

First-Order Logic

FOL Overview

- First Order logic (FOL) is a powerful knowledge representation (KR) system
- Used in AI systems in various ways, e.g., to
 - Directly represent & reason about concepts & objects
 - Formally specify meaning of KR systems (e.g., [OWL](#))
 - Form programming languages (e.g., [Prolog](#)) and [rule-based systems](#)
 - Make semantic database systems ([Datalog](#)) and Knowledge graphs ([Wikidata](#))
 - Provide features useful in neural network deep learning systems

First-order logic

- First-order logic (FOL) models the world in terms of
 - **Objects**, which are things with individual identities
 - **Properties** of objects that distinguish them from others
 - **Relations** that hold among sets of objects
 - **Functions**, a subset of relations where there is only one “value” for any given “input”
- Examples:
 - Objects: students, lectures, companies, cars ...
 - Relations: isa, hasBrother, biggerThan, outside, hasPart, color, occursAfter, owns, visits, precedes, ...
 - Properties: blue, oval, even, large, ...
 - Functions: hasFather, hasSSN, ...

User provides

- **Constant symbols** representing individuals in world
 - BarackObama, Green, John, 3, “John Smith”
- **Predicate symbols** map individuals to truth values
 - greater(5,3)
 - green(Grass)
 - color(Grass, Green)
 - hasBrother(John, Robert)
- **Function symbols** map individuals to individuals
 - hasFather(SashaObama) = BarackObama
 - colorOf(Sky) = Blue

What do these mean?

- User should also indicate what these mean in a way that humans will understand
 - i.e., map to their own internal representations
- May be done via a combination of
 - Choosing good names for formal terms, e.g. calling a concept HumanBeing instead of [Q5](#)
 - Comments in the definition `#human being`
 - Descriptions and examples in documentation
 - Reference to other representations , e.g., sameAs [/m/0dgw95](#) in Freebase and [Person](#) in schema.org
 - Give examples like *Donald Trump* and *Luke Skywalker* to help distinguish the concepts of a real and fictional person

FOL Provides

- **Variable symbols**
 - e.g., X , Y , $?foo$, $?number$
- **Connectives**
 - Same as propositional logic: not (\neg), and (\wedge), or (\vee), implies (\rightarrow), iff (\leftrightarrow), equivalence (\equiv), ...
- **Quantifiers**
 - Universal $\forall x$ or **(Ax)**
 - Existential $\exists x$ or **(Ex)**

Sentences: built from terms and atoms

- **term** (denoting an individual): constant or variable symbol, or n-place function of n terms, e.g.:
 - Constants: john, umbc
 - Variables: X, Y, Z
 - Functions: mother_of(john), phone(mother(x))
- **Ground terms** have no variables in them
 - **Ground:** john, father_of(father_of(john))
 - **Not Ground:** father_of(X)
- Syntax may vary: maybe variables must start with a “?” or a capital letter

Sentences: built from terms and atoms


- **atomic sentences** (which are either true or false) are n-place predicates of n terms, e.g.:
 - green(kermit)
 - between(philadelphia, baltimore, dc)
 - loves(X, mother(X))
- **complex sentences** formed from atomic ones connected by the standard logical connectives with quantifiers if there are variables, e.g.:
 - loves(mary, john) \vee loves(mary, bill)
 - $\forall x$ loves(mary, x)

What do atomic sentences mean?

- Unary predicates typically encode a **type**
 - `muppet(Kermit)`: kermit is a kind of muppet
 - `green(kermit)`: kermit is a kind of green thing
 - `integer(X)`: x is a kind of integer
- Non-unary predicates typically encode relations or properties
 - `Loves(john, mary)`
 - `Greater_than(2, 1)`
 - `Between(newYork, philadelphia, baltimore)`
 - `hasName(john, "John Smith")`

Ontology

- Designing a logic representation is like designing a model in an object-oriented language
- **Ontology:** a “formal naming and definition of the types, properties and relations of entities for a domain of discourse”
- E.g.: schema.org ontology used to put semantic data on Web pages to help search engines
 - Here’s the [semantic markup](#) Google sees about [CSEE department of UMBC](#).



Sentences: built from terms and atoms

- **quantified sentences** adds quantifiers \forall and \exists
 $\forall x \text{ loves}(x, \text{mother}(x))$
 $\exists x \text{ number}(x) \wedge \text{greater}(x, 100), \text{prime}(x)$
- **well-formed formula (wff)**: a sentence with no *free* variables or where all variables are *bound* by a universal or existential *quantifier*
In $(\forall x)P(x, y)$ x is bound & y is free so it's not a wff

Quantifiers: \forall and \exists

- **Universal quantification**
 - $(\forall x)P(X)$ means P holds for **all** values of X in the domain associated with variable¹
 - E.g., $(\forall X) \text{dolphin}(X) \rightarrow \text{mammal}(X)$
- **Existential quantification**
 - $(\exists x)P(X)$ means P holds for **some** value of X in domain associated with variable
 - E.g., $(\exists X) \text{mammal}(X) \wedge \text{lays_eggs}(X)$
 - This lets us make statements about an object without identifying it

¹ a variable's domain is often not explicitly stated and is assumed by the context

Universal Quantifier: \forall

- **Universal quantifiers typically used with *implies* to form *rules*:**

Logic: $\forall X \text{ student}(X) \rightarrow \text{smart}(X)$

Means: All students are smart

- **Universal quantification *rarely* used without implies:**

Logic: $\forall X \text{ student}(X) \wedge \text{smart}(X)$

Means: Everything is a student and is smart

- **What about this, though:**

—*Logic:* $\forall X \text{ alive}(X) \vee \text{dead}(X)$

—Means: everything is either alive or dead

Universal Quantifier: \forall

- **What about this, though:**
 - *Logic:* $\forall X \text{ alive}(X) \vee \text{dead}(X)$
 - Means: everything is either alive or dead
- **Can be rewritten using a standard tautology**
 - $A \vee B \equiv \sim A \rightarrow B$
- **Giving both of these (since $A \vee B \equiv B \vee A$)**
 - $\forall X \sim \text{alive}(X) \rightarrow \text{dead}(X)$
 - $\forall X \sim \text{dead}(X) \rightarrow \text{alive}(X)$

Existential Quantifier: \exists

- Existential quantifiers usually used with **and** to specify a list of properties about an individual

Logic: $(\exists X) \text{ student}(X) \wedge \text{ smart}(X)$

Meaning: There is a student who is smart

- Common mistake: represent this in FOL as:

Logic: $(\exists X) \text{ student}(X) \rightarrow \text{ smart}(X)$

Meaning: ?

Existential Quantifier: \exists

- Existential quantifiers usually used with **and** to specify a list of properties about an individual

Logic: $(\exists X) \text{ student}(X) \wedge \text{smart}(X)$

Meaning: There is a student who is smart

- Common mistake: represent this in FOL as:

Logic: $(\exists X) \text{ student}(X) \rightarrow \text{smart}(X)$

$P \rightarrow Q = \sim P \vee Q$

$\exists X \text{ student}(X) \rightarrow \text{smart}(X) = \exists X \sim \text{student}(X) \vee \text{smart}(X)$

Meaning: There's something that is either not a student or is smart

Quantifier Scope

- FOL sentences have structure, like programs
- In particular, variables in a sentence have a **scope**
- Suppose we want to say “everyone who is alive loves someone”
 $(\forall X) \text{ alive}(X) \rightarrow (\exists Y) \text{ loves}(X, Y)$
- Here’s how we scope the variables

$(\forall X) \text{ alive}(X) \rightarrow (\exists Y) \text{ loves}(X, Y)$

 Scope of x
 Scope of y

Quantifier Scope

- **Switching order of two universal quantifiers *does not* change the meaning**
 - $(\forall X)(\forall Y)P(X,Y) \leftrightarrow (\forall Y)(\forall X) P(X,Y)$
 - Dogs hate cats (i.e., all dogs hate all cats)
- **You can switch order of existential quantifiers**
 - $(\exists X)(\exists Y)P(X,Y) \leftrightarrow (\exists Y)(\exists X) P(X,Y)$
 - A cat killed a dog
- **Switching order of universal and existential quantifiers *does* change meaning:**
 - Everyone likes someone: $(\forall X)(\exists Y) \text{ likes}(X,Y)$
 - Someone is liked by everyone: $(\exists Y)(\forall X) \text{ likes}(X,Y)$

```
def verify1():
```

```
    # Everyone likes someone:  $(\forall x)(\exists y) \text{ likes}(x,y)$ 
```

```
    for p1 in people():
```

```
        foundLike = False
```

```
        for p2 in people():
```

```
            if likes(p1, p2):
```

```
                foundLike = True
```

```
                break
```

```
        if not foundLike:
```

```
            print(p1, 'does not like anyone ☹')
```

```
            return False
```

```
    return True
```

Every person has at least one individual that they like.

Procedural example 1

```
def verify2():
```

```
    # Someone is liked by everyone:  $(\exists y)(\forall x) \text{ likes}(x,y)$ 
```

```
    for p2 in people():
```

```
        foundHater = False
```

```
        for p1 in people():
```

```
            if not likes(p1, p2):
```

```
                foundHater = True
```

```
                break
```

```
        if not foundHater
```

```
            print(p2, 'is liked by everyone 😊')
```

```
            return True
```

```
    return False
```

There is a person who is liked by every person in the universe.

Procedural example 2

Connections between \forall and \exists

- We can relate sentences involving \forall and \exists using extensions to De Morgan's laws:
 1. $(\forall x) P(x) \leftrightarrow \neg(\exists x) \neg P(x)$
 2. $\neg(\forall x) P(x) \leftrightarrow (\exists x) \neg P(x)$
 3. $(\exists x) P(x) \leftrightarrow \neg (\forall x) \neg P(x)$
 4. $\neg(\exists x) P(x) \leftrightarrow (\forall x) \neg P(x)$
- Examples
 1. All dogs sleep \leftrightarrow There is no dog that doesn't sleep
 2. Not all dogs bark \leftrightarrow There is a dog that doesn't bark
 3. There is a dog that talks \leftrightarrow Not all dogs can't talk
 4. No dog likes cats \leftrightarrow All dogs don't like cats

Translating English to FOL

Every gardener likes the sun

All purple mushrooms are poisonous

Translating English to FOL

Every gardener likes the sun

$\forall x \text{ gardener}(x) \rightarrow \text{likes}(x, \text{Sun})$

All purple mushrooms are poisonous

Translating English to FOL

Every gardener likes the sun

$$\forall x \text{ gardener}(x) \rightarrow \text{likes}(x, \text{Sun})$$

All purple mushrooms are poisonous

$$\forall x (\text{mushroom}(x) \wedge \text{purple}(x)) \rightarrow \text{poisonous}(x)$$

Translating English to FOL

Every gardener likes the sun

$$\forall x \text{ gardener}(x) \rightarrow \text{likes}(x, \text{Sun})$$

All purple mushrooms are poisonous

$$\forall x (\text{mushroom}(x) \wedge \text{purple}(x)) \rightarrow \text{poisonous}(x)$$

No purple mushroom is poisonous (two ways)

Translating English to FOL

Every gardener likes the sun

$$\forall x \text{ gardener}(x) \rightarrow \text{likes}(x, \text{Sun})$$

All purple mushrooms are poisonous

$$\forall x (\text{mushroom}(x) \wedge \text{purple}(x)) \rightarrow \text{poisonous}(x)$$

No purple mushroom is poisonous (two ways)

$$\neg \exists x \text{ purple}(x) \wedge \text{mushroom}(x) \wedge \text{poisonous}(x)$$

Translating English to FOL

Every gardener likes the sun

$$\forall x \text{ gardener}(x) \rightarrow \text{likes}(x, \text{Sun})$$

All purple mushrooms are poisonous

$$\forall x (\text{mushroom}(x) \wedge \text{purple}(x)) \rightarrow \text{poisonous}(x)$$

No purple mushroom is poisonous (two ways)

$$\neg \exists x \text{ purple}(x) \wedge \text{mushroom}(x) \wedge \text{poisonous}(x)$$

$$\forall x (\text{mushroom}(x) \wedge \text{purple}(x)) \rightarrow \neg \text{poisonous}(x)$$

English to FOL: Counting



Use = predicate to identify different individuals

- There are at least two purple mushrooms

$$\exists x \exists y \text{ mushroom}(x) \wedge \text{purple}(x) \wedge \text{mushroom}(y) \wedge \text{purple}(y) \wedge \neg(x=y)$$

- There are exactly two purple mushrooms

$$\begin{aligned} &\exists x \exists y \text{ mushroom}(x) \wedge \text{purple}(x) \wedge \text{mushroom}(y) \wedge \text{purple}(y) \wedge \\ &\neg(x=y) \wedge \\ &\forall z (\text{mushroom}(z) \wedge \text{purple}(z)) \rightarrow ((x=z) \vee (y=z)) \end{aligned}$$

Saying there are 802 different Pokemon is hard!

Direct use of FOL is not for everything!

Translating English to FOL



What do these mean?

- You can fool *some of* the people *all of* the time
- You can fool *all of* the people *some of* the time

Translating English to FOL



What do these mean?

Both English statements are ambiguous

- **You can fool *some of the people* *all of the time***

There is a nonempty subset of people so easily fooled that you can fool that subset every time.

For any given time, there is a non-empty subset at that time that you can fool

- **You can fool *all of the people* *some of the time***

There are one or more times when it's possible to fool everyone.

Each individual can be fooled at some point in time

Translating English to FOL



You can fool *some of* the people *all of* the time

There is a nonempty group of people so easily fooled that you can fool that group every time

≡ There's (at least) one person you can fool every time

$\exists x \forall t \text{ person}(x) \wedge \text{time}(t) \rightarrow \text{canFool}(x, t)$

For any given time, there is a non-empty group at that time that you can fool

≡ For every time, there's a person at that time that you can fool

$\forall t \exists x \text{ person}(x) \wedge \text{time}(t) \rightarrow \text{canFool}(x, t)$

Translating English to FOL



You can fool *all of* the people *some of* the time

There's at least one time when you can fool everyone

$$\exists t \forall x \text{time}(t) \wedge \text{person}(x) \rightarrow \text{canFool}(x, t)$$

Everybody can be fooled at some point in time


$$\forall x \exists t \text{person}(x) \wedge \text{time}(t) \rightarrow \text{canFool}(x, t)$$

Representation Design



- Many options for representing even a simple fact, e.g., something's color as red, green or blue, e.g.:
 - `green(kermit)`
 - `color(kermit, green)`
 - `hasProperty(kermit, color, green)`
- Choice can influence how easy it is to use
- Last option of representing properties & relations as [triples](#) used by modern [knowledge graphs](#)
 - Easy to ask: What color is Kermit? What are Kermit's properties?, What green things are there? Tell me everything you know, ...

Simple genealogy KB in FOL



Design a knowledge base using FOL that

- Has facts of immediate family relations, e.g., spouses, parents, etc.
- Defines more complex relations (ancestors, relatives)
- Detect conflicts, e.g., you are your own parent
- Infers relations, e.g., grandparent from parent
- Answers queries about relationships between people

Extend with relations and constraints

- Simple two argument relations, e.g.
 - spouse, has_child, has_parent
- Add general constraints to relations, e.g.
 - $\text{spouse}(X,Y) \Rightarrow \sim (X = Y)$
 - $\text{spouse}(X,Y) \Rightarrow \text{person}(X) \wedge \text{person}(Y)$
- Add FOL sentences for inference, e.g.
 - $\text{spouse}(X,Y) \Leftrightarrow \text{spouse}(Y,X)$
- Add instance data
 - e.g., $\text{spouse}(\text{Djt}, \text{Mt})$

Example: A simple genealogy KB in FOL

Predicates:

- parent(x, y), child(x, y), father(x, y), daughter(x, y), etc.
- spouse(x, y), husband(x, y), wife(x, y)
- ancestor(x, y), descendant(x, y)
- male(x), female(y)
- relative(x, y)

Facts (ground terms):

- husband(Joe, Mary), son(Fred, Joe)
- spouse(John, Nancy), male(John), son(Mark, Nancy)
- father(Jack, Nancy), daughter(Linda, Jack)
- daughter(Liz, Linda)
- etc.

Example Axioms

$(\forall x, y) \text{ parent}(x, y) \leftrightarrow \text{child}(y, x)$

$(\forall x, y) \text{ father}(x, y) \leftrightarrow \text{parent}(x, y) \wedge \text{male}(x)$

$(\forall x, y) \text{ mother}(x, y) \leftrightarrow \text{parent}(x, y) \wedge \text{female}(x)$


$(\forall x, y) \text{ daughter}(x, y) \leftrightarrow \text{child}(x, y) \wedge \text{female}(x)$

$(\forall x, y) \text{ son}(x, y) \leftrightarrow \text{child}(x, y) \wedge \text{male}(x)$

$(\forall x, y) \text{ husband}(x, y) \leftrightarrow \text{spouse}(x, y) \wedge \text{male}(x)$

$(\forall x, y) \text{ spouse}(x, y) \leftrightarrow \text{spouse}(y, x)$

...



Axioms, definitions and theorems

- **Axioms**: facts and rules that capture (important) facts & concepts in a domain; used to prove **theorems**
 - Dependent axioms can make reasoning faster, however
 - Choosing a good set of axioms is a design problem
- A **definition** of a predicate is of the form “ $p(X) \leftrightarrow \dots$ ” and can be decomposed into two parts
 - **Necessary** description: “ $p(x) \rightarrow \dots$ ”
 - **Sufficient** description “ $p(x) \leftarrow \dots$ ”
 - Some concepts have definitions (e.g., triangle) and some don't (e.g., person)

More on definitions

Example: define $\text{father}(x, y)$ by $\text{parent}(x, y)$ & $\text{male}(x)$

- **$\text{parent}(x, y)$** is a necessary (but not sufficient) description of $\text{father}(x, y)$

$$\text{father}(x, y) \rightarrow \text{parent}(x, y)$$

- **$\text{parent}(x, y) \wedge \text{male}(x) \wedge \text{age}(x, 35)$** is a sufficient (but not necessary) description of $\text{father}(x, y)$:

$$\text{father}(x, y) \leftarrow \text{parent}(x, y) \wedge \text{male}(x) \wedge \text{age}(x, 35)$$

- **$\text{parent}(x, y) \wedge \text{male}(x)$** is a necessary and sufficient description of $\text{father}(x, y)$

$$\text{parent}(x, y) \wedge \text{male}(x) \leftrightarrow \text{father}(x, y)$$

Higher-order logic

- FOL only lets us quantify over variables, and **variables can only range over objects**
- HOL allows us to quantify over relations, e.g.

“two functions are equal iff they produce the same value for all arguments”

$$\forall f \forall g (f = g) \leftrightarrow (\forall x f(x) = g(x))$$

- E.g.: (quantify over predicates)

$$\forall r \text{ transitive}(r) \rightarrow (\forall xyz) r(x,y) \wedge r(y,z) \rightarrow r(x,z)$$

- More expressive, but reasoning is undecide-able, in general

Examples of FOL in use



- Semantics of W3C's [Semantic Web](#) stack (RDF, RDFS, OWL) is defined in FOL
- [OWL](#) Full is equivalent to FOL
- Other OWL profiles support a subset of FOL and are more efficient
- FOL oriented knowledge representation systems have many user friendly tools
- E.g.: Protégé for creating, editing and exploring OWL ontologies

Examples of FOL in use

Many practical approaches embrace the approach that “some data is better than none”

- The semantics of schema.org is only defined in natural language text
- [Wikidata](https://www.wikidata.org/)’s knowledge graph has a rich schema
 - Many constraint/logical violations are flagged with warnings
 - However, not all, see this Wikidata query that finds people who are their own mother or father



Automated inference for FOL

- Automated inference for FOL is harder than PL
 - Variables can take on an infinite number of possible values from their domains
 - Hence there are potentially an infinite number of ways to apply the Universal Elimination rule
- Godel's Completeness Theorem says that FOL entailment is only semi-decidable
 - If a sentence is **true** given a set of axioms, there is a procedure that will determine this
 - If a sentence is **false**, there's no guarantee a procedure will ever discover this — it **may never halt**


Generalized Modus Ponens (GMP)

- Modus Ponens: $P, P \Rightarrow Q \models Q$
- Generalized Modus Ponens extends this to rules in FOL
- Combines And-Introduction, Universal-Elimination, and Modus Ponens, e.g.
 - given $P(c), Q(c), \forall x P(x) \wedge Q(x) \rightarrow R(x)$
 - derive $R(c)$
- Must deal with
 - more than one condition on rule's left side
 - variables

Forward & Backward Reasoning

- We often talk about two reasoning strategies:
 - Forward chaining and
 - Backward chaining
- Both are equally powerful, but optimized for different use cases
- You can also have a mixed strategy

Forward chaining




- Proofs start with given axioms/premises in KB, deriving new sentences using GMP until the goal/query sentence is derived
 - The process follows a chain of rules and facts going from the KB to the conclusion
- This defines a **forward-chaining** inference procedure because it moves “forward” from the KB to the goal [eventually]
- Inference using GMP is **sound** and **complete** for KBs containing **only Horn clauses**

Forward chaining example

- KB:
 - $\text{allergies}(X) \rightarrow \text{sneeze}(X)$
 - $\text{cat}(Y) \wedge \text{allergicToCats}(X) \rightarrow \text{allergies}(X)$
 - $\text{cat}(\text{felix})$
 - $\text{allergicToCats}(\text{mary})$
- Goal:
 - $\text{sneeze}(\text{mary})$

Backward chaining



- **Backward-chaining** deduction using GMP is also **complete** for KBs containing **only Horn clauses**
- Proofs start with the goal query, find rules with that conclusion, and then tries to prove each of the antecedents in the rule
- Keep going until you reach premises
- Avoid loops by checking if new subgoal is already on the goal stack
- Avoid repeated work: use a cache to check if new subgoal already proved true or failed

Backward chaining example

- KB:
 - $\text{allergies}(X) \rightarrow \text{sneeze}(X)$
 - $\text{cat}(Y) \wedge \text{allergicToCats}(X) \rightarrow \text{allergies}(X)$
 - $\text{cat}(\text{felix})$
 - $\text{allergicToCats}(\text{mary})$
- Goal:
 - $\text{sneeze}(\text{mary})$

Forward vs. backward chaining

- Forward chaining is data-driven
 - Automatic, unconscious processing, e.g., object recognition, routine decisions
 - May do lots of work that is irrelevant to the goal
 - Efficient when you want to compute all conclusions
- Backward chaining is goal-driven, better for problem-solving and query answering
 - Where are my keys? How do I get to my next class?
 - Complexity can be much less than linear w.r.t KB size
 - Efficient when you want one or a few decisions
 - Good where the underlying facts are changing

Mixed strategy

- Many practical reasoning systems do both forward and backward chaining
- The way you encode a rule determines how it is used, as in
 - % this is a forward chaining rule
`spouse(X,Y) => spouse(Y,X).`
 - % this is a backward chaining rule
`wife(X,Y) <= spouse(X,Y), female(X).`
- Given a model of the rules you have and the kind of reason you need to do, it's possible to decide which to encode as FC and which as BC rules.

Summary

- Logical agents apply inference to KB to derive new information and make decisions
- Basic concepts of logic:
 - Syntax: formal structure of sentences
 - Semantics: truth of sentences wrt models
 - Entailment: necessary truth of one sentence given another
 - Inference: deriving sentences from other sentences
 - Soundness: derivations produce only entailed sentences
 - Completeness: derivations produce all entailed sentences
- FC and BC linear time, complete for Horn clauses
- Resolution is sound and complete for propositional and first-order logic